

# Clustering (2/2)

**Clustering: advanced aspects**

**DASH: Data Science e Análise Não Supervisionada**

Rui Henriques, [rmch@tecnico.ulisboa.pt](mailto:rmch@tecnico.ulisboa.pt)

Instituto Superior Técnico, Universidade de Lisboa

# Outline

- Distances: advanced notes
- Clustering approaches
  - $k$ -means
  - model-based (EM)
  - deep learning
- Evaluation recap
- Advanced aspects


# Outline

- **Distances: advanced notes**
- Clustering approaches
  - $k$ -means
  - model-based (EM)
  - deep learning
- Evaluation recap
- Advanced aspects

# Recall: clustering

**Cluster:** group of observations

**Cluster analysis:** group observations into clusters according to their (dis)similarity:  
observations in the same cluster are more similar than those in different clusters



Main ingredients for clustering: ***distance + approach***

# Clustering: distance and approach

**Distance metrics** depend on the:


- type of input variables:

- *numeric* distances, e.g. Euclidean,  $d(\mathbf{a}, \mathbf{b}) = \sqrt{\sum_{j=1}^m (a_j - b_j)^2}$

- *nominal* distances, e.g. Hamming
- *ordinal* encodings, how?
- *non-iid attributes*, how?

- data structures



- (multivariate) time series
- image
- text
- others: spatiotemporal, events, relational...



# Categorical features

How do we represent categorical features in a multivariate space?

- Let us start with **binary variables** (e.g. gender)
  - a binary feature can be mapped into one axis in a  $m$ -dimensional space where values are constrained to 0 or 1
- Ordinal variables** (e.g. high-moderate-low)
  - we can find a numeric encoding  $f$  for an ordinal variable
    - e.g.  $f(\text{high}) = 5, f(\text{moderate}) = 1, f(\text{low}) = 0$
  - in the absence of an encoding, one can assume equally spaced integers (e.g. 0, 1, 2...). Problems?
- Let us finally consider **nominal variables** with cardinality higher than 2 (e.g. Europe, Africa, America)
  - challenge? As there is no order, we cannot position values along a single axis/dimension
  - solution? Create  $p$  axes for a nominal variable with cardinality  $p$ 
    - only one of the  $p$  axes is active (1), while remaining are inactive (0)
    - this operation is called **dummification**
    - some machine learning methods depend on dummification



# Mixed data

- A data space can be composed by **non-iid numeric variables**
  - *i.i.d.* stands for **independent** and **identically distributed**
  - two numeric variables may be correlated, i.e. not independent
  - two numeric variables may have distinct distributions (e.g.  $U(0,100)$  and  $N(3,5)$ ), i.e. non-identically distributed
- Observations can have numeric, nominal and ordinal features
  - we informally term such data as **mixed data**
  - examples?



# Distances for mixed data

- First concern: two non-identically distributed variables – e.g.  $Y_1 \sim U(0,1)$  and  $Y_2 \sim U(0,100)$ 
  - *Problem?* In the given example, distances are most affected by variable  $Y_2$ 
    - yet why should this variable have higher weight than  $Y_1$ !
  - *Possible solution?* Normalize variables
    - in the given example,  $Y_2 \sim U'(0,1)$  using for instance min-max scaling
- Second concern: how to deal with simultaneous categoric and numeric variables?
  - distance can be a composition:  $d(\mathbf{x}_1, \mathbf{x}_2) = \alpha d_{numeric}(\mathbf{x}_1, \mathbf{x}_2) + \beta d_{categoric}(\mathbf{x}_1, \mathbf{x}_2)$   
where  $\alpha$  and  $\beta$  are parameters that reveal the relevance of each component, generally  $\alpha + \beta = 1$ 
    - parameters can be fixed based on the number of variables or available domain knowledge



# Handling complex data structures

Two major paradigms

- dedicated distances for series/image/text/events (see previous class!)
- feature extraction to map complex data structures into tabular data
  - *manual extraction* often incomplete and error-prone
  - *automated extraction* susceptible to diverse idiosyncrasies
    - examples: statistics on a sliding window for *time series*, structure–texture–color cues for *images*, term frequencies for *text*, geometric features for *trajectories*...
  - **problem:** *feature extraction* and subsequent *engineering* entail 90% of overall ML effort!
  - **solution:** **learn representations**
    - “*A good representation is one that makes subsequent learning task easier*”  
Goodfellow et al. 2016 (including clustering!)


# Numeric encodings of complex data

- A good **representation** is:
  - compact (*minimal*)
  - explanatory (*sufficient*)
  - disentangled (*independent factors*)
  - informative (*make subsequent problem solving easier*)
- How to learn numeric representations from complex data?
  - **neural networks** ⇒ check our class on **representation learning!**
- Implications?
  - classic distances can be applied
  - yet given the above properties: **cosine distance** recommended



# Distances vs correlation



- Consider the scenario where you have three documents (observations):  $d_1$ ,  $d_2$  and  $d_3$ 
  - each document is characterized by the frequency of terms in the text (features)
- Now consider we enter a query  $q$  to retrieve documents similar to  $q$
- Euclidean distance is... a bad idea as it is **large** for vectors of **different lengths**
  - the Euclidean distance between  $q$  and  $d_2$  is large even though the distribution of terms in the query  $q$  and  $d_2$  are very similar



# Vector norm


The normalization of an observation, should not be mislead with *variable normalization*

- variable normalization guarantees that the measurements of a given variable yield statistical properties of interest, e.g. [0,1] or  $N(0,1)$



# From angles to cosines


- the following two notions are equivalent:
  - *distance* of the angle between two vectors
  - *similarity* of the cosine between two vectors
- cosine is a monotonically decreasing function for the interval  $[0^\circ, 180^\circ]$



# Cosine similarity

$$\cos(\mathbf{x}_1, \mathbf{x}_2) = \frac{\mathbf{x}_1 \cdot \mathbf{x}_2}{\|\mathbf{x}_1\| \|\mathbf{x}_2\|} = \frac{\sum_i x_{1i} x_{2i}}{\sqrt{\sum_i x_{1i}^2} \sqrt{\sum_i x_{2i}^2}}$$


- $x_{1i}$  is the feature from variable  $y_i$  in observation  $\mathbf{x}_1$
- $\|\mathbf{x}_1\|$  and  $\|\mathbf{x}_2\|$  are 2-norm lengths of vectors  $\mathbf{x}_1$  and  $\mathbf{x}_2$
- $\cos(\mathbf{x}_1, \mathbf{x}_2)$  is the (cosine) similarity of  $\mathbf{x}_1$  and  $\mathbf{x}_2$



For length-normalized vectors, cosine similarity is simply the dot product (or scalar product):

$$\cos(\mathbf{x}_1, \mathbf{x}_2) = \mathbf{x}_1 \cdot \mathbf{x}_2 = \sum_{i=1}^m x_{1i} x_{2i}$$

# Cosine similarity



**Euclidean distance**

$$l_2(\mathbf{x}_1, \mathbf{x}_2) = \sqrt{|2 - 6|^2 + |8 - 3|^2} = \sqrt{41}$$

**Manhattan distance**

$$l_1(\mathbf{x}_1, \mathbf{x}_2) = |2 - 6| + |8 - 3| = 9$$

**Cosine similarity**

$$\cos(\mathbf{x}_1, \mathbf{x}_2) = \frac{\mathbf{x}_1 \cdot \mathbf{x}_2}{\|\mathbf{x}_1\| \|\mathbf{x}_2\|} = \frac{2 \times 6 + 8 \times 3}{\sqrt{2^2 + 8^2} \sqrt{6^2 + 3^2}} = 0.65$$

# Distances in clustering

- Distances (or similarities) applied between:
  - two observations  $d(\mathbf{x}_i, \mathbf{x}_j)$
  - one observation and a cluster  $d(\mathbf{x}_i, \mathbf{c}_j)$
  - two clusters  $d(\mathbf{c}_i, \mathbf{c}_j)$
- **Cluster:**  $\mathbf{c}_j = \{\mathbf{x} \mid d(\mathbf{x}, \mathbf{c}_j) = \min_i d(\mathbf{x}, \mathbf{c}_i)\}$
- **Centroid** of a cluster as its mass center:  $\bar{\mathbf{c}}_j$  (e.g., mean/mode value per real/categorical variable)
- Squared **error** of clustering solution:  $E = \sum_{k=1}^K \sum_{\mathbf{x} \in \mathbf{c}_k} d(\mathbf{x}, \bar{\mathbf{c}}_k)^2$

# Outline

- Distances: advanced notes
- **Clustering approaches**
  - ***k*-means**
  - **model-based (EM)**
  - **deep learning**
- Evaluation recap
- Advanced aspects

# Approaches

## Partitioning ⇐

- Create partitions and iteratively update them (e.g.  $k$ -means,  $k$ -modes,  $k$ -medoids)

## Hierarchical

- Create hierarchical decomposition of data points

## Density-based

- Group points based on connectivity and density functions

## Model-based ⇐

- Data are seen as a mixture of distributions (e.g. EM)

## Deep learning ⇐

- End-to-end neural networks for high-efficacy clustering

# Partitioning algorithms


Given  $k$  clusters:

1. partition observations into  $k$  non-empty subsets
2. compute the centroid  $c_j$  of each subset
  - centroid is the center of mass: e.g. mean or median centers
3. reassign each observation to the cluster with the nearest centroid
4. goto step 2, stop when:
  - i) assignment does not change, ii)  $|E^{new} - E^{old}| < \varepsilon$ , or iii) max iterations is reached

## Variants

- centroid calculus
- selection of the initial seeds
- adjustments for **batches** of observations (instead of all) for very large datasets


# *k*-means




$k=2$

Arbitrarily choose  
 $k$  object as initial  
cluster center

Assign each objects to most similar center




Update the cluster means




reassign

Update the cluster means



# *k*-medoids

- **numeric** data: centroid as the mean or median
- **categoric** data: centroid as the mode – *k*-modes [Huang'98]
  - frequency-based procedure to update modes of clusters
- **mixed** data: centroid combining mean and modes (***k*-prototype**)
- ***k*-medoids**: the **most centrally located observation** in a cluster is the centroid
  - observation with minimum average distance to all observations in the cluster
- What is the algorithm most robust to outliers: *k*-means or *k*-medoids?





# *k*-means: challenges



- **Efficiency:**  $O(tkn)$   $n=\#$ observation,  $k=\#$ clusters,  $t=\#$ iterations, usually  $k, t \ll n$
- **Problems**
  - dependent on initialization
  - sensitive to outliers
  - sensitive noisy data
    - noise can substantially distort centroids
  - not suitable to discover clusters with *non-convex shapes*
    - deal only with clusters with spherical symmetrical point distribution
  - need to specify  $k$ , the *number* of clusters, in advance
  - convergence?

# Limitations of $k$ -means

Different sizes




Different densities




# Limitations of $k$ -means





Non-globular shapes



Fixed number of clusters





# Overcoming $k$ -means limitations




# Model-based vs partitioning approaches

- $k$ -means algorithm performs a **hard assignment** of data points to clusters, in which each data point is associated uniquely with one cluster
  - in  $k$ -means the shape of the cluster is described by Euclidean distance function
- Model-based clustering makes a **soft assignment** based on the posterior probabilities
  - Expectation Maximization (EM) is the paradigmatic algorithm



# EM: example




# EM: principles

- Fix the number of clusters, assume each cluster follows a specific distribution
  - multivariate Gaussian mixture for numeric variables
  - frequentist view for discrete variables
  - joint probabilities under independence assumption for mixed data
- EM algorithm
  - Iterate between two steps until convergence


**Expectation step:** assign observations to clusters

**Maximation step:** update the model parameters (adjust distributions)




# EM and soft clustering

- each observation has a likelihood of being generated by every cluster,  $p(\mathbf{x}|c)$ 
  - can be assigned to the cluster with higher likelihood,  $\text{argmax}_c p(c|\mathbf{x})$
  - can be assigned to more than one cluster (if both satisfy minimum likelihood) or none
- tackles major problems of density-based approaches (e.g. DBSCAN cannot partition data below)



# Deep learning approaches

- Check our presentation on *representation learning* to recover the foundations!
- Two major paradigms
  1. **learn good representations** from data (whether simple or complex data structures) using deep learning *followed by classic clustering stances* (e.g.  $k$ -means)
    - addresses data idiosyncrasies (e.g., non-iid variables)



2. apply emerging **end-to-end deep learning** pipelines for clustering

# Outline

- Distances: advanced notes
- Clustering approaches
  - $k$ -means
  - model-based (EM)
  - deep learning
- **Evaluation recap**
- Advanced aspects

# Recall: clustering evaluation



- 3 kinds of measures: external, internal and relative indexes
- **External** (supervised): extent to which cluster labels match true labels
  - requires prior or expert knowledge
- **Internal** (unsupervised): goodness without external information
  - how well they are separated (e.g. silhouette)
  - should be independent from algorithm-specific functions (unbiased)
- **Relative**: compare different cluster analyses (different parameters/algorithms)

# Statistical significance

- **Why?** need to robustly interpret measure
  - if our measure of evaluation has the value, 10, is that good, fair, or poor?  $\Rightarrow p$ -value
- **How?** Compare the values of an index that result from a randomized dataset against the index values from the target cluster analysis
  - If the value of the index is unlikely, then results are significant
  - To test superiority between two clustering solutions, simpler statistical tests can be considered
  - Interesting fact: the more “atypical” a clustering result is, the more likely it represents valid structure in the data




# Statistical significance

- Compare SSE of 0.005 against SSE of clusters in 500 sets of random data points of size 100 (histogram)



# Statistical significance

- Are correlations of -0.92 and -0.58 statistically significant?  
Histogram shows correlation on randomized data with K-means




# Outline

- Distances: advanced notes
- Clustering approaches
  - $k$ -means
  - model-based (EM)
  - deep learning
- Evaluation recap
- **Advanced aspects**

# Recall: clustering modes

- memberships: **hard** (observation belongs to a single cluster) *versus* **soft** (each observation has a probability of belonging to each cluster – e.g. fuzzy, model-based clustering)
- supervision
  - **unsupervised** (*default*): cluster observations without knowing their outcome variable
  - **semi-supervised**: cluster observations when labels of some observations may be known *or* pairs of observations are known to belong to the same cluster
  - **“supervised”**: cluster observations in the presence of variables of interest (e.g. targets as inputs, class-conditional clustering)
- separation of clusters: **exclusive** versus **non-exclusive** (overlapping clusters)
- **complete** versus **partial** (cluster some objects)
- **uniform** versus **weighted attributes** or **observations**



# Visualizing clustering solutions


**Key principles** for knowledge discovery (to be mastered during our course!)

- Compute the variables that are better separated by the clusters
  - add a clusters column to the dataset and run ANOVA (*f-classif* in *sklearn*) to assess the **discriminative power** of each **input variable** (use p-values to rank them by importance)
- Retrieve the **centroids**
- Compute the **cluster-conditional distributions** for the most important features
- Retrieve **observation memberships/distances** to each cluster/centroid
- **Visualize** clustering solutions in a **2D or 3D space**
  - select specific features of interest (importance or domain knowledge)
  - project the original m-dimensional space into a 2D or 3D space using uMAP, PCA, tSNE

# Visualizing clustering solutions

**Key principles** for knowledge discovery (to be mastered during our course!)

- Retrieve the **centroids** and/or **medoids**
- Compute the **cluster-conditional distributions** for the most relevant variables




- Retrieve **observation memberships/distances** to each cluster/centroid







# Visualizing clustering solutions

**Key principles** for knowledge discovery (to be mastered during our course!)

- Most informative variables? The ones that were better separated by the clusters
  - add a column with clusters to the dataset
  - run ANOVA (*f-classif* in *sklearn*) to assess the **discriminative power** of each **input variable**
  - use p-values to rank variables by importance
- **Visualize** clustering solutions in a **2D or 3D space**
  - select specific features of interest (importance or domain knowledge)
  - project the original m-dimensional space into a 2D or 3D space using uMAP, PCA, tSNE




# Importance of seeding initial centroids





# Importance of seeding initial centroids

- If there are  $k$  real clusters...
  - ... chance of selecting one centroid from each cluster is small when  $k$  is large
    - e.g. if  $k = 10$ , then probability =  $10!/10^{10} = 0.00036$
- **Difficulty:** sometimes centroids readjust in 'right' way, sometimes don't
- **Solution**
  - multiple runs (helps, but probability is not on your side)
  - hierarchical clustering to determine initial centroids
  - select more than  $k$  initial centroids and then select among these initial centroids (select most widely separated)
  - postprocessing


# Number of clusters?





How many clusters?




Six Clusters



Two Clusters




Four Clusters




# Number of clusters?


$k = 1$ : cohesion (SSE) is 873.0



$k = 2$ : cohesion (SSE) is 173.1



$k = 3$ : cohesion (SSE) is 133.6



# Number of clusters?

## Knee/elbow method:

- plot the SSE for different  $k$
- find the elbow. Example: abrupt changes ( $k=2$ ) are highly suggestive of two clusters in the data

## Other methods

- plot metrics less sensitive to the #clusters (e.g., silhouette)
- find the  $k$  with maximum value or higher neighborhood average centered on  $k$



# Pre- and post-processing

- **Pre-processing**

- normalize data
- data reduction and transformation
- remove outliers

- **Post-processing**

- eliminate small clusters that may represent outliers
- split '*loose*' clusters (clusters with relatively high SSE)
- merge clusters that are '*close*' (clusters with relatively low SSE)
- these steps can be integrated within the clustering process

# Requirements

- clustering **quality**
- ability to deal with different **data types**
  - attribute: numeric, nominal, ordinal and mixtures
  - structure: vectors, events, time series, etc.
  - dimensionality
- ability to deal with **noise**, **missings** and **outliers**
- discovery of clusters with arbitrary **shape**
- **actionability**
- **efficiency/scalability**
- **interpretability** and usability
- **parameter-free** (unfixed number of clusters)
- clustering in the presence of domain knowledge and user-specified constraints

# How data characteristics affect clustering?

- sparseness
- variable domain
- data structure
- data distribution
- size and dimensionality
- noise and outliers



# Thank you!

Rui Henriques

[rmch@tecnico.ulisboa.pt](mailto:rmch@tecnico.ulisboa.pt)