

# Dimensionality reduction

Feature selection, principal component analysis, discriminant analysis

DASH: Data Science e Análise Não Supervisionada

Rui Henriques, [rmch@tecnico.ulisboa.pt](mailto:rmch@tecnico.ulisboa.pt)

Instituto Superior Técnico, Universidade de Lisboa

# Outline

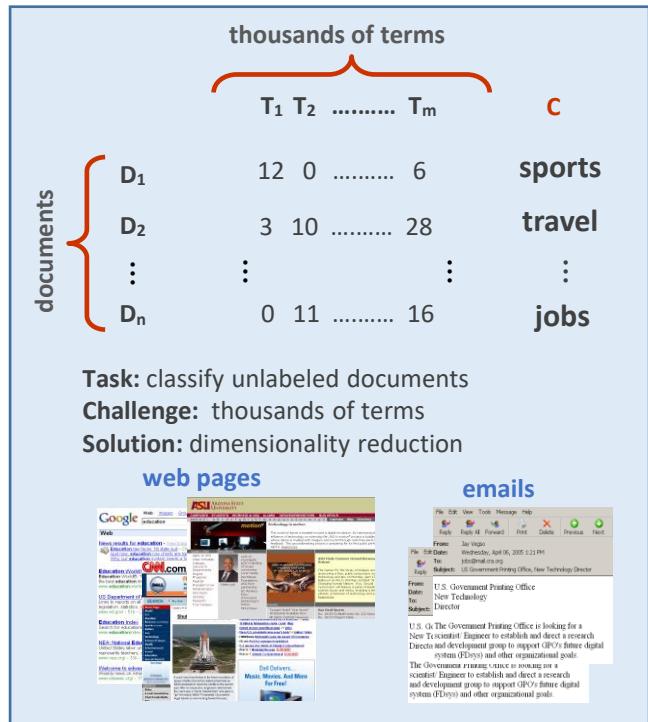
- Dimensionality reduction: why and how
- Feature selection
- Linear transformations
  - algebra ground
  - principal component analysis
  - compression and reconstruction
  - non-linear kernels
  - linear discriminant analysis
- Subspace selection
- Evaluation
- Data reduction

# Motivation

- At a first glimpse, increasing the number of variables should improve learning...
- In practice, including more variables can degrade performance (i.e. **curse of dimensionality**)
  - *challenges*: learning complexity and generalization difficulty (over/underfitting)
  - common definition of **high-dimensionality**:  $|Y| \gg |X|$  (i.e.  $m \gg n$ )
- The number of training observations required increases **exponentially** with dimensionality
- How then can we learn in high-dimensional data spaces with a limited number of observations?
  - **dimensionality reduction**

# Data domains with high-dimensionality

- **text** and **web content** data (*left*)
- **social** behavioral data
- **biological data**
  - genetic variants (millions per individual)
  - gene expression (>20k genes)
  - molecular concentrations (metabolites, proteins...)
- **healthcare** data (clinical records)
- **consumer data**
- **signal, audio, image** and **video** data

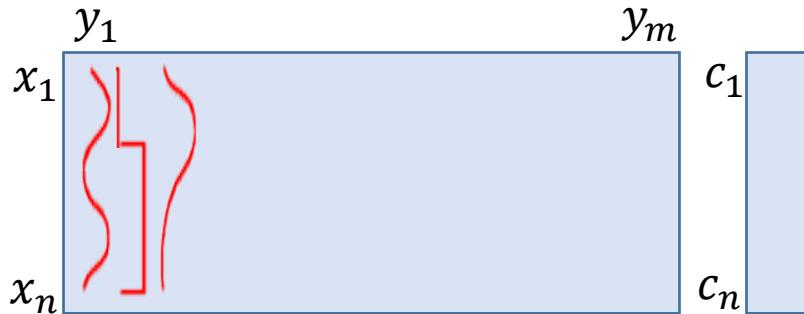


# Challenges

High-dimensional data analysis:

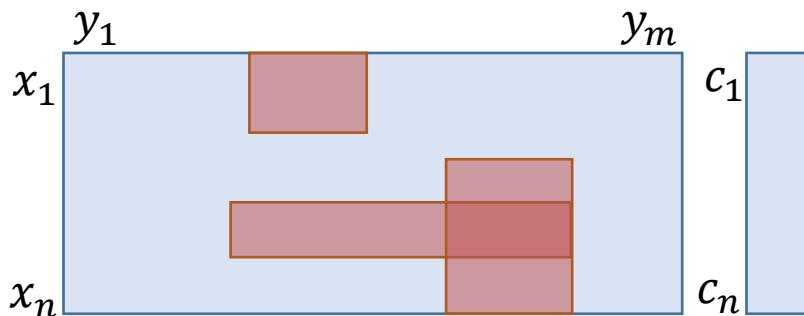
- learning **complexity**: large data space
- **generalization difficulty** (insufficient observations)
  - **overfitting** risk
  - **underfitting** risk
- how to evaluate these risks?
  - behavior for varying data size
  - compare training and testing error
  - variability of errors across testing folds
  - bias and variance components of error

# Generalization: overfitting and underfitting risks



## Overfitting

- inability to discard non-informative and/or non-discriminative data



## Underfitting

- exclusion of informative or discriminative data from the learning

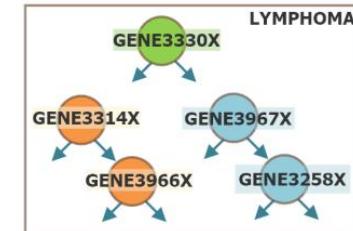
# Global and local learning

## "Global" learning

- descriptors (e.g., *clusters*) and predictors (e.g., *discriminants*)
  - all features considered, often equally relevant in joint probabilistic stances (e.g. *naïve Bayes*)
- Problem? **overfitting**
  - what if only a few variables are relevant?

## "Local" learning

- descriptors (*patterns*) and predictors (e.g. *decision trees*, *kNN*)
  - few combined features or observations as long as they are informative or discriminative
- Problem? **underfitting**
  - many potentially relevant features or observations neglected



# Dimensionality reduction

Some **goals** of dimensionality reduction:

- Guide **supervised learning** (focus on discriminative regions)
- Guide **unsupervised learning** (focus on informative regions)
- **Visualization** (project high-dim data into interpretable low-dim data)
- **Data compression** (efficient storage and retrieval)
- **Noise removal** (denoising data)
- **Speed-up** learning
- **Describe** the underlying properties of data
- Guarantee simplicity and comprehensibility of mined results
- Map **multimedia data** (image and signal data) into feature-based data
- Support matrix operations (inverse, rank determination, approximation)...



# How?

## 1. Feature selection

## 2. Feature extraction/transformation

- principal component analysis
- linear discriminant analysis
- representation learning

## 3. Sparse kernels and regularization in parametric models to exclude non-relevant parameters

- neural networks, support vector machines, discriminant analysis...

## 4. Subspace selection to jointly select variables and observations

- pattern mining, decision trees and random forests
- associative classifiers and decision tables

# Dimensionality reduction

- Project the  $m$ -dimensional observations into a  $k$ -dimensional space ( $k \ll m$ )
  - preserve most of relevant information or structure from data
- Solve the learning problem in low dimensions
- Two major approaches
  - **feature selection**
    - choosing a subset of all features
$$[y_1, y_2, \dots, y_m] \rightarrow [y_{i1}, y_{i2}, \dots, y_{ik}]$$
  - **feature extraction**
    - creating new features by combining existing ones
$$[y_1, y_2, \dots, y_m] \rightarrow [d_1, d_2, \dots, d_k] \text{ using } f([y_{i1}, y_{i2}, \dots, y_{im}])$$

# Outline

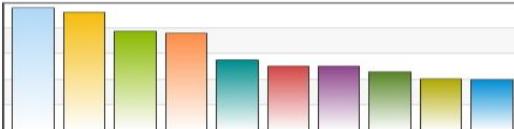
- Dimensionality reduction: why and how
- **Feature selection: recall**
- Linear transformations
  - algebra ground
  - principal component analysis
  - compression and reconstruction
  - non-linear kernels
  - linear discriminant analysis
- Subspace selection
- Evaluation
- Data reduction

# Feature selection

- Optimal subset of features according to an objective function
  - differs from feature extraction where output features are (non-)linear combinations of original features
- **Objective criteria**
  - **unsupervised** setting: *variables with higher variability and entropy*
  - **supervised** setting: *maximize discrimination/correlation with target variable*
- **Perspectives**
  - subset search vs. feature ranking
  - models: filter vs. wrapper

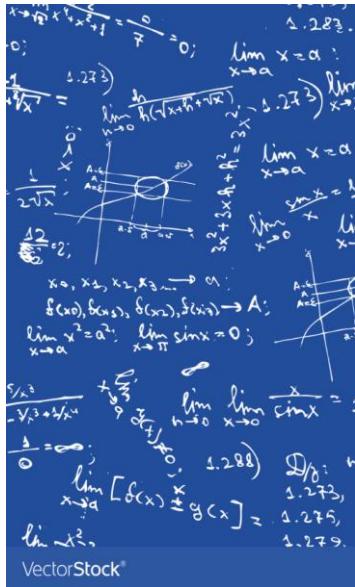
# Feature ranking

- Weighting each individual feature according to objective criteria (e.g. information gain)



- Sort and select **top-ranked features**: a) threshold, b)  $k$ -top, or c) percentile
- Disadvantages
  - hard to determine threshold or  $k$
  - unable to consider correlation between features
- Advantages
  - efficient  $O(m)$  no need to test combination of features (subset optimality)

# Measures for feature ranking



Goodness of a feature/feature subset:

- **information measures**
- **correlation/association measures**
- **distance measures**
- **accuracy measures**

# Information measures in classification

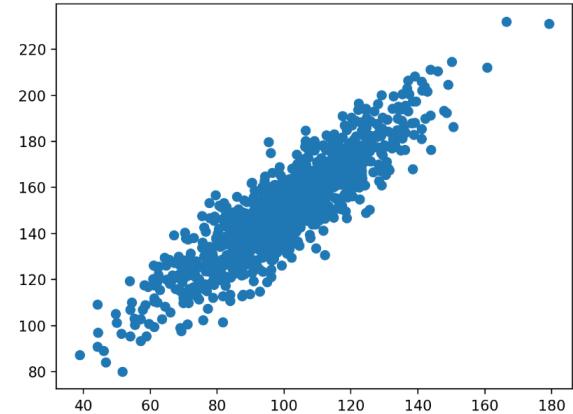
- **$\chi^2$  test** (chi2 in Python)
  - robust for input categorical variables
  - two values:  $\chi^2$  statistic (higher the better),  $p$ -value

```
iris = datasets.load_iris()  
X, y = iris.data, iris.target  
  
chi2, pval = chi2(X, y)  
  
[ 10.81782088  3.59449902  116.16984746  67.24482759]  
[ 4.47651e-03  1.6575e-01  5.943443e-26  2.50017e-15]
```

- **ANOVA test** (f\_classif in Python)
  - robust for numeric input variables
  - also valid for categorical input with numerical output
  - two values: F-value statistic (higher the better),  $p$ -value

# Information measures in regression

- **correlation** for numeric input variables:
  - Pearson and Spearman (see previous lectures)
  - $F$ -statistic (higher the better) and accompanying  $p$ -value
- **mutual information** or **ANOVA** for categorical input variables
- want to abstract the type of input variables?  
use `f_regression` from `sklearn`



# Information measures in description

- as a **filter**: measure feature importance and select top- $k$  features or above threshold
  - **unsupervised** setting: **high entropy**
  - supervised setting: next slides
- as a **wrapper**: assess learning performance with varying subsets of features
  - simple: measure feature importance and test descriptors on top- $k$  features with varying  $k$
  - advanced: assess descriptors with different subsets of variables (irrespectively of top- $k$ )
  - how to assess descriptors?
    - e.g. clustering quality – any problem when considering silhouette?

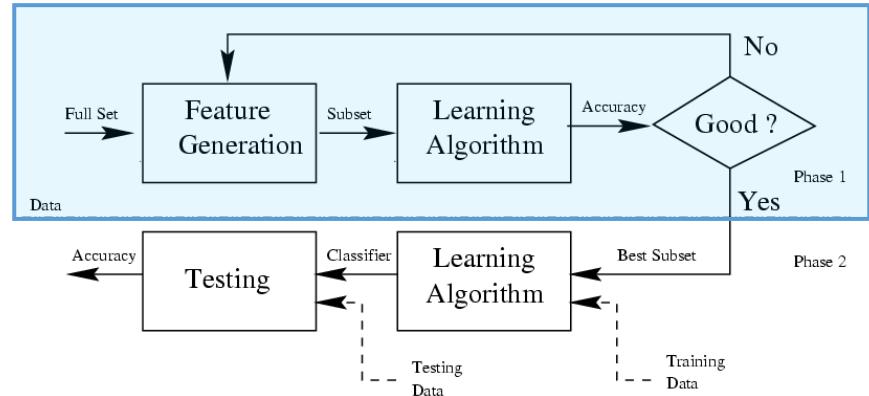
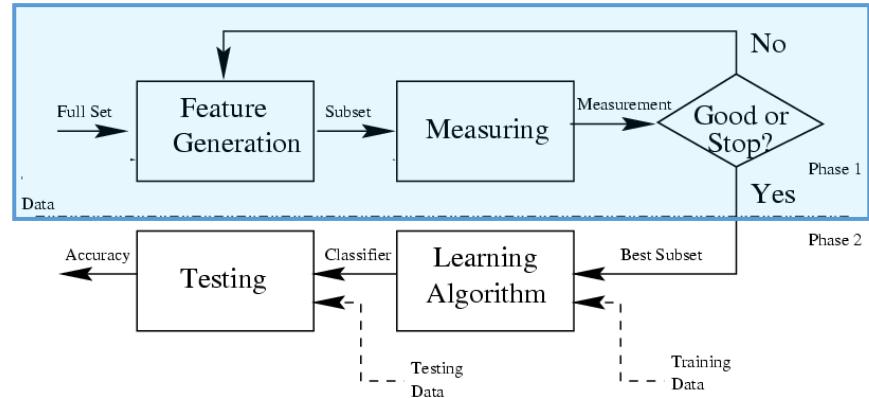
# Feature selection: filter vs. wrapper

## Filter model

- independent from the learning algorithm
  - efficient and no learning biases
- rely on general characteristics of data

## Wrapper model

- evaluated on a given descriptor/predictor
- descriptive utility or predictive accuracy as a goodness measure
- better yet dependent on the learning approach
- more computationally expensive

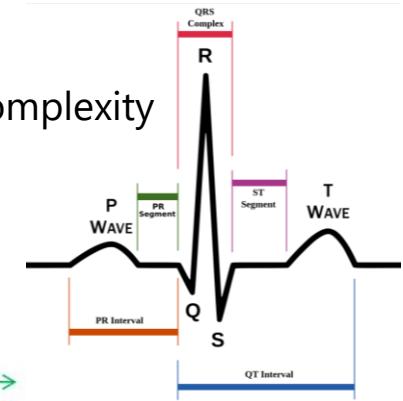
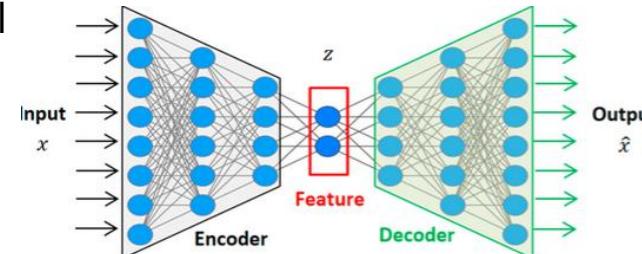


# Outline

- Dimensionality reduction: why and how
- Feature selection
- **Linear transformations**
  - **algebra ground**
  - **principal component analysis**
  - compression and reconstruction
  - non-linear kernels
  - linear discriminant analysis
- Subspace selection
- Evaluation
- Data reduction

# Feature extraction

- Feature extraction used to
  - **extract relevant features from complex data** to handle structural complexity
    - extract domain-specific features of interest
      - e.g. PQRST statistics from ECG signal
    - extract domain-independent statistics of interest
      - e.g. spectral statistics from time series and images
  - learn latent features using neural networks (**data representation**)



- **learn transformations** that explain simple multivariate data with reduced dimensionality

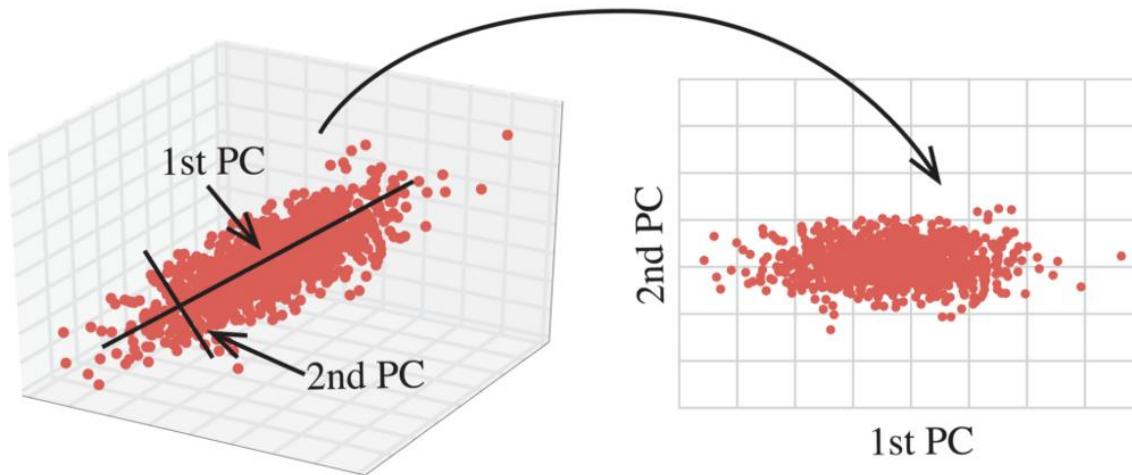
# Feature extraction

Feature extraction used to

- extract relevant features from complex data to handle structural complexity
- **learn transformations** that explain simple multivariate data with reduced dimensionality: **TODAY!**
  - classical approaches aim at finding a *linear transformation*
    - Why? simple to compute and analytically tractable
    - Goal: reduction that preserves as much information in data as possible
      - paradigmatic case: **Principal Component Analysis** (PCA)
  - simple extensions available
    - *non-linear transformations* (kernel trick)
    - accommodate discriminative power in *supervised settings*
      - Goal: reduction that best separates the data
        - paradigmatic case: **Linear Discriminant Analysis** (LDA)



# Algebra ground for PCA



Axes of greater variance given by *eigenvectors of covariance matrix*

# Algebra ground

- A **covariance matrix** measures the tendency of two variables to vary in the same direction
  - positive (negative) covariance denote positive (inverse) correlation
    - nevertheless magnitude of values not easily interpretable
    - when normalizing covariance by their variances we obtain linear correlation in [-1,1]
  - covariance matrix is symmetric and positive-definite
- Remember
  - sample covariance:  $n - 1$  in the denominator (Bessel's correction)

$$cov(y_1, y_2) = \frac{\sum_{i=1}^n (a_{1i} - \bar{y}_1) \cdot (a_{2i} - \bar{y}_2)}{n - 1}$$

- whole population:  $n$  is the denominator

$$cov(y_1, y_2) = \frac{\sum_{i=1}^n (a_{1i} - \bar{y}_1) \cdot (a_{2i} - \bar{y}_2)}{n}$$

# Algebra ground

- **Covariance matrix**, given  $m$  attributes,  $y_1, \dots, y_m$

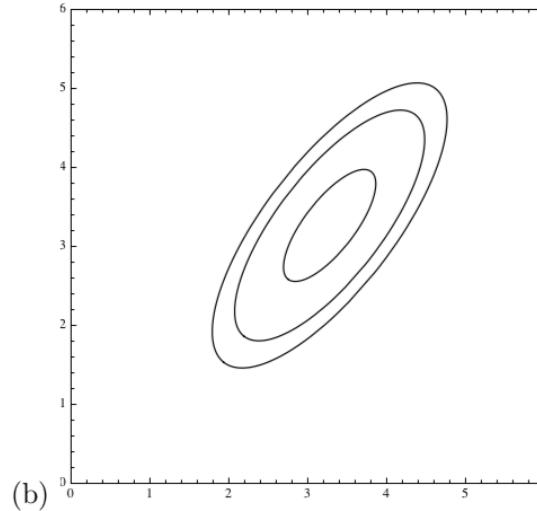
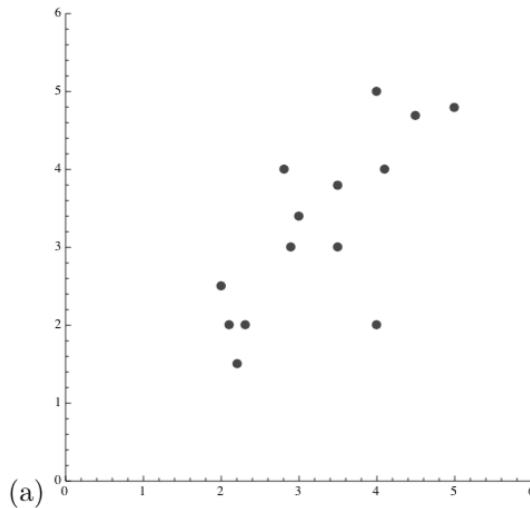
$$C^{m \times m} = (c_{ij} | i,j = 1..m), \text{ where } c_{ij} = \text{cov}(y_i, y_j)$$

Data	Hours(H)	Mark(M)
9	39	
15	56	
25	93	
14	61	
10	50	
18	75	
0	32	
16	85	
5	42	
19	70	
16	66	
20	80	
Totals	167	749
Averages	13.92	62.42

$$\begin{aligned} C^{m \times m} &= \begin{pmatrix} \text{cov}(H, H) & \text{cov}(H, M) \\ \text{cov}(M, H) & \text{cov}(M, M) \end{pmatrix} \\ &= \begin{pmatrix} \text{var}(H) & 104.5 \\ 104.5 & \text{var}(M) \end{pmatrix} \\ &= \begin{pmatrix} 47.7 & 104.5 \\ 104.5 & 370 \end{pmatrix} \end{aligned}$$

# Algebra ground

- The covariance matrix of the data points defines the ellipses of equiprobability on the right



# Eigenvalues and eigenvectors

- Let  $C$  be a  $m \times m$  covariance matrix
- Vectors  $\mathbf{v}$  having same direction as  $C\mathbf{v}$  are called eigenvectors
  - eigenvectors define the linear composition of attributes
- In the equation  $C\mathbf{v} = \lambda\mathbf{v}$ ,  $\lambda$  is called an eigenvalue of  $C$
- Example:

$$\begin{pmatrix} 2 & 3 \\ 2 & 1 \end{pmatrix} \begin{pmatrix} 3 \\ 2 \end{pmatrix} = \begin{pmatrix} 12 \\ 8 \end{pmatrix} = 4 \begin{pmatrix} 3 \\ 2 \end{pmatrix}$$

$$\mathbf{v} = [3 \ 2]^T \text{ and } \lambda = 4$$

meaning that data is described by  $y_{new} = 3y_1 + 2y_2$

# Eigenvalues and eigenvectors

- $A\mathbf{v} = \lambda\mathbf{v} \Leftrightarrow (A - \lambda I)\mathbf{v} = 0$
- Given  $A$ , how to calculate  $\mathbf{v}$  and  $\lambda$ :
  - determine roots to  $\det(A - \lambda I) = 0$ , roots are eigenvalues  $\lambda$
  - solve  $(A - \lambda I)\mathbf{v} = 0$  for each  $\lambda$  to obtain eigenvectors  $\mathbf{v}$

$y_1$	$y_2$
-5.1	9.25
14.9	20.25
5.9	33.25
5.9	-30.75
...	...
-9.1	-10.75
-9.1	-21.75
5.9	19.25

$$C = \begin{pmatrix} 2 & 0.8 \\ 0.8 & 0.6 \end{pmatrix}$$

Eigenvectors and eigenvalues:

$$\mathbf{v}_1 = [0.91, 0.41], \lambda_1 = 2.36$$

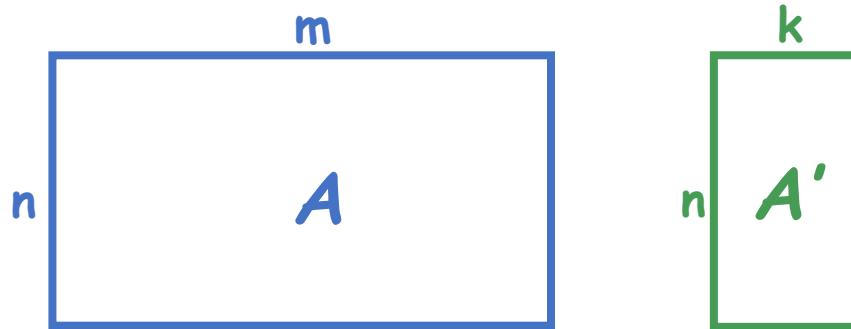
$$\mathbf{v}_2 = [-0.41, 0.91], \lambda_2 = 0.23$$

$$\mathbf{x}_i = (0.91 \quad 0.41) \begin{pmatrix} a_{i1} \\ a_{i2} \end{pmatrix}$$

$y_{new}$
-0.8
21.9
19
-7.2
...
-12.7
-17.2
13.3

# Dimensionality reduction

- Map data with  $m$  variables into  $k$  variables (such that  $k < m$ ) without significant loss



- Residual variation:* information in  $A$  not retained in  $A'$
- Trade-off: dimensionality ( $k$ ) and interpretability *versus* information loss

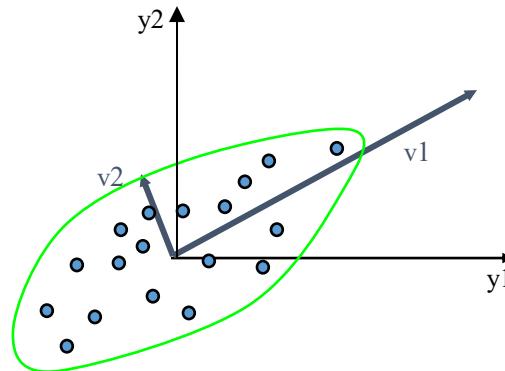
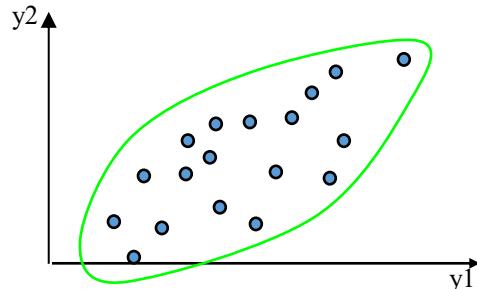
# Linear transformations

## Intuition:

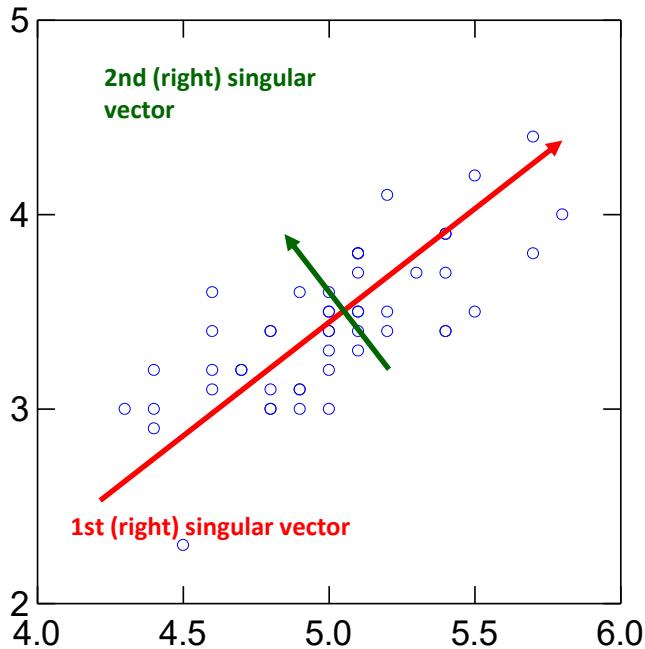
- find the axis that shows the greatest variation
- project all points into this axis

## How:

- move origin to the center of the dataset
- find the eigenvectors and eigenvalues of the data covariance matrix
  - the eigenvectors define the new data space



# Singular value decomposition (SVD)



**1<sup>st</sup> singular vector:** direction of maximal variance

$\lambda_1$ : how much data variance is explained by 1<sup>st</sup> vector

**2<sup>nd</sup> singular vector:** direction of maximal variance, after removing projection of 1<sup>st</sup> vector

$\lambda_2$ : how much data variance is explained by the 2<sup>nd</sup> vector

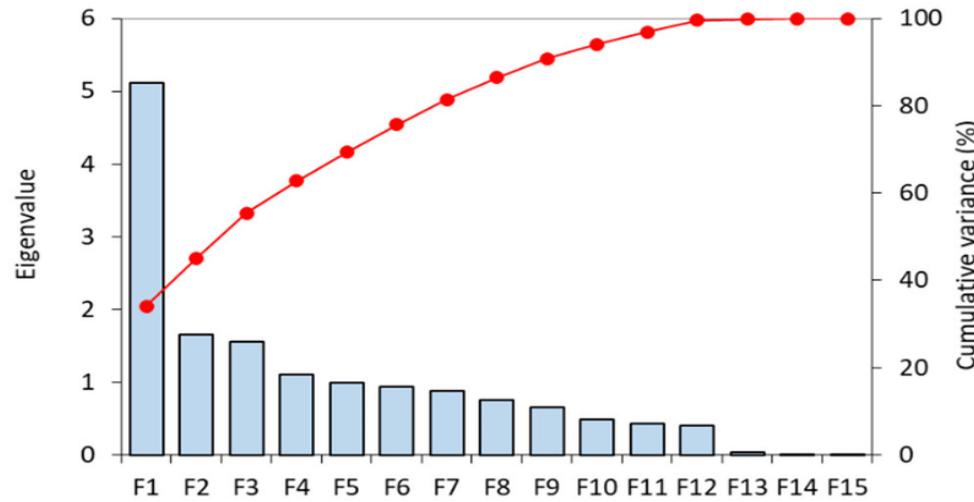
...

**$k^{\text{th}}$  singular vector ...**

(until unexplained variance below threshold)

# Principal component analysis (PCA)

- PCA is SVD done on *centered* data
  - singular vector/value = eigenvector/value
- First component (PC1): highest eigenvalue (direction with greatest variation)
- Second component (PC2): direction with maximum variation orthogonal to PC1
- ...



# Component selection

- The variance in the direction of the  $k^{\text{th}}$  eigenvector is given by the eigenvalue  $\lambda_k$
- Singular values can be used to estimate how many components to keep
- **Rule of thumb:** keep enough to explain 85% of the variation

$$\frac{\sum_{j=1}^k \lambda_j}{\sum_{j=1}^m \lambda_j} \approx 0.85 \quad \begin{array}{l} \text{if } k = m, \text{ we preserve 100\%} \\ \text{of the original variation} \end{array}$$

- depending on the reduced dimensionality and learning needs: 90%, 95%, 98% also common
- **Karhunen-Loeve** (KL) transform is PCA without subsequent removal of components
  - no information loss

# Principal component analysis

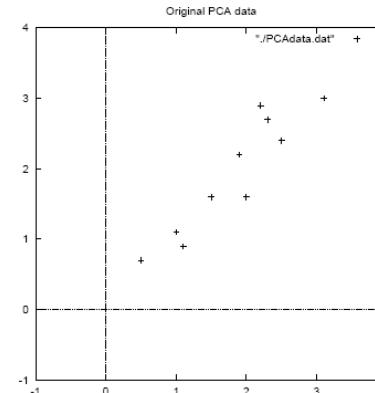
- PCA projects data along the directions where the data varies most
  - directions are determined by the eigenvectors with the largest eigenvalues
  - reduction can imply information loss, yet PCA preserves as much information as possible
- Components (summary variables)
  - linear combinations of the original variables
  - uncorrelated with each other
  - the largest eigenvalues are called ***principal components***
    - the eigenvalue is the magnitude of eigenvector, defining the direction's variance
    - in general: *only few components needed to capture most data variability*

# Principal component analysis

- Revising how

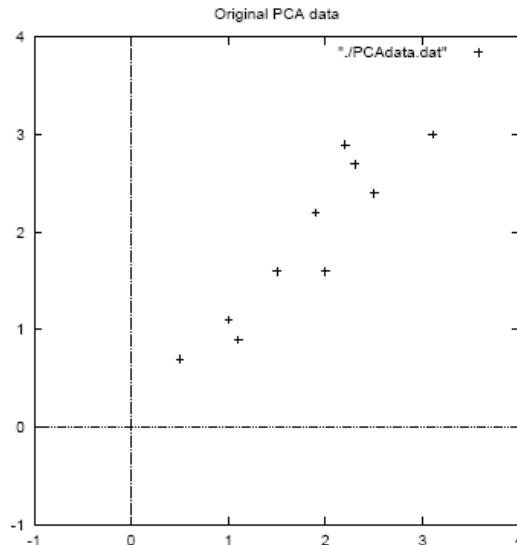
1. Subtract mean of each variable
2. Compute the covariance matrix  $m \times m$  (*scatter of data*)
3. Compute eigenvalues,  $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_m$ , and eigenvectors,  $v_1, v_2, \dots, v_m$
4. Keep the large  $k$  eigenvalues ( $k \leq m$ ) and construct the transformed space
5. Transform the dataset  $D \rightarrow D'$

- **Exercise:** apply PCA on the following dataset



# PCA: Example

## 1. Centering data



$y_1$	$y_2$	$y_1$	$y_2$
2.5	2.4	.69	.49
0.5	0.7	-1.31	-1.21
2.2	2.9	.39	.99
1.9	2.2	.09	.29
3.1	3.0	OriginalData =	CenteredData =
2.3	2.7	1.29	1.09
2	1.6	.49	.79
1	1.1	.19	-.31
1.5	1.6	-.81	-.81
1.1	0.9	-.31	-.31
		-.71	-1.01
mean 1.81   1.91		mean 0   0	

# PCA example

2. Calculate the covariance matrix:

$$cov = \begin{pmatrix} & y_1 & y_2 \\ .616555556 & & .615444444 \\ .615444444 & & .716555556 \end{pmatrix}_{y_1 \ y_2}$$

3. Calculate its (unit) eigenvectors and eigenvalues

$$eigenvalues = \begin{pmatrix} .0490833989 \\ 1.28402771 \end{pmatrix} \quad eigenvectors = \begin{pmatrix} -.735178656 & -.677873399 \\ .677873399 & -.735178656 \end{pmatrix}$$

4. Order eigenvectors by eigenvalue, highest to lowest and select top  $p$

$$\mathbf{v}_1 = \begin{pmatrix} -.6779 \\ -.7352 \end{pmatrix} \quad \lambda_1 = 1.284 \quad \mathbf{v}_2 = \begin{pmatrix} -.7352 \\ .6779 \end{pmatrix} \quad \lambda_2 = .0491$$

... and construct the transformed feature vector

$$FeatureVector(k=2) = \begin{pmatrix} -.6779 & -.7352 \\ -.7352 & .6779 \end{pmatrix} \quad FeatureVector(k=1) = \begin{pmatrix} -.6779 \\ -.7352 \end{pmatrix}$$

# PCA example

## 5. Derive the new data set

$$\text{TransformedData} = \text{RowFeatureVector} \times \text{RowDataAdjust}$$

$$\text{DataAdjusted} = \begin{pmatrix} .69 & -1.31 & .39 & .09 & 1.29 & .49 & .19 & -.81 & -.31 & -.71 \\ .49 & -1.21 & .99 & .29 & 1.09 & .79 & -.31 & -.81 & -.31 & -1.01 \end{pmatrix}$$

$$\text{FeatureVector}(p = 2)^T = \begin{pmatrix} -.6779 & -.7352 \\ -.7352 & .6779 \end{pmatrix}$$

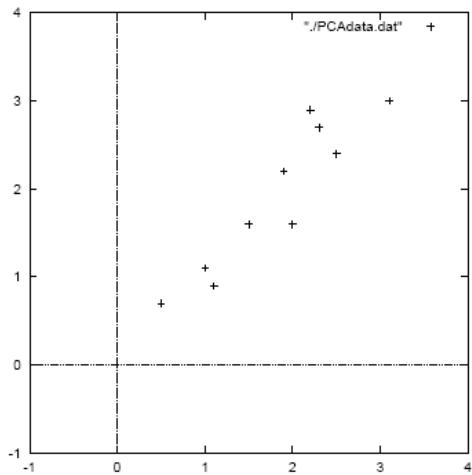
$$\text{FeatureVector}(p = 1)^T = (-.6779 \quad -.7352)$$

	$c_1$	$c_2$
	-.827970186	-.175115307
	1.77758033	.142857227
	-.992197494	.384374989
	-.274210416	.130417207
TransformedData (p=2) =	-1.67580142	-.209498461
	-.912949103	.175282444
	.0991094375	-.349824698
	1.14457216	.0464172582
	.438046137	.0177646297
	1.22382056	-.162675287

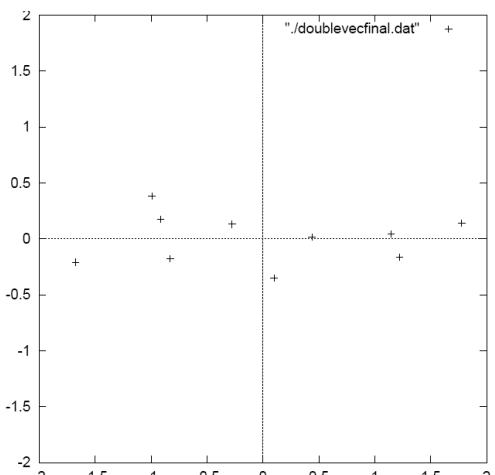
	$c_1$
	-.827970186
	1.77758033
	-.992197494
	-.274210416
TransformedData (p=1) =	-1.67580142
	-.912949103
	.0991094375
	1.14457216
	.438046137
	1.22382056

# PCA example

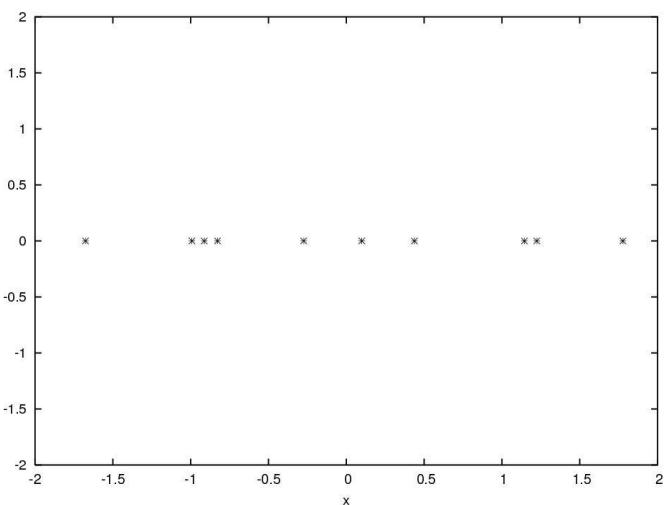
Original data



Data transformed with  
 $k = 2$  eigenvectors



Data transformed with  
 $k = 1$  eigenvector



# Outline

- Dimensionality reduction: why and how
- Feature selection
- **Linear transformations**
  - algebra ground
  - principal component analysis
  - **compression and reconstruction**
  - non-linear kernels
  - linear discriminant analysis
- Subspace selection
- Evaluation
- Data reduction

# Compression

- **Goal:** compress data without information loss
- Why? Remember...
  - **description**
    - compactness, interpretability, simplicity
    - transforms can be used to explain relevant information (variable dependencies)
    - how? Assess which variables are contributing more to each component
      - remember a component is a linear composition of variables
  - guide **prediction** on the compressed data space
  - **denoising** data
  - **efficient** data storage, retrieval, learning

# Recovering original data

- The original vectors can be reconstructed using its principal components
- Given the reduced data space, how to retrieve/decompress old data?

$$\text{DataRecovered} = (\text{FeatureVector} \times \text{TransformedData}) + \text{OriginalMean}$$

we did:  $\text{TransformedData} = \text{FeatureVector}^T \times \text{DataAdjust}^T$

so we can do:  $\text{DataAdjust}^T = (\text{FeatureVector}^T)^{-1} \times \text{TransformedData}$

(orthogonal property)  $= \text{FeatureVector} \times \text{TransformedData}$

and:  $\text{DataRecovered} = \text{DataAdjust} + \text{OriginalMean}$

# Reconstruction: example

$$\text{FeatureVector}(p = 2)^T = \begin{pmatrix} -.6779 & -.7352 \\ -.7352 & .6779 \end{pmatrix}$$

$$\text{FeatureVector}(p = 1)^T = (-.6779 \quad -.7352)$$

$y_1$		$c_1$		$c_2$		using both components		after uncentering	
						$y'_1$	$y'_2$	$y'_1$	$y'_2$
2.5	2.4	-.827970186	-.175115307	0.69	0.49	2.5	2.4		
0.5	0.7	1.77758033	.142857227	-1.31	-1.21	0.5	0.7		
2.2	2.9	-.992197494	.384374989	0.39	0.99	2.2	2.9		
1.9	2.2	-.274210416	.130417207	0.09	0.29	1.9	2.2		
3.1	3.0	-1.67580142	-.209498461	1.29	1.09	3.1	3		
2.3	2.7	-.912949103	.175282444	0.49	0.79	2.3	2.7		
2	1.6	.0991094375	-.349824698	0.19	-0.31	2	1.6		
1	1.1	1.14457216	.0464172582	-0.81	-0.81	1	1.1		
1.5	1.6	.438046137	.0177646297	-0.31	-0.31	1.5	1.6		
1.1	0.9	1.22382056	-.162675287	-0.71	-1.01	1.1	0.9		

$$\text{DataRecovered} = (\text{FeatureVector} \times \text{TransformedData}) + \text{OriginalMean}$$

# Reconstruction error

- PCA minimizes the reconstruction error:  $\|\mathbf{x} - \hat{\mathbf{x}}\|$
- It can be shown that the reconstruction error is:  $error = 1/2 \sum_{i=k+1}^m \lambda_i$ 
  - using 2 components: recovery error = 0 (from previous slide)
  - using 1 component

$y'_1$	$y'_2$	$y^*_1$	$y^*_2$
0.56	0.61	2.4	2.5
-1.20	-1.31	0.6	0.6
0.67	0.73	2.5	2.6
0.19	0.20	2.0	2.1
1.14	1.23	2.9	3.1
0.62	0.67	2.4	2.6
-0.07	-0.07	1.7	1.8
-0.78	-0.84	1.0	1.1
-0.30	-0.32	1.5	1.6
-0.83	-0.90	1.0	1.0

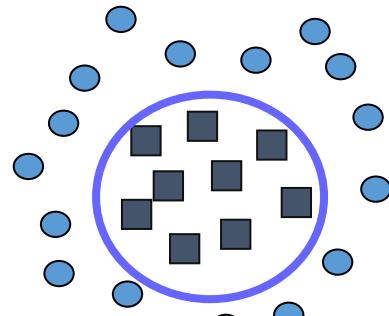
$$error = 0.245 = \frac{0.49}{2} = \frac{\lambda}{2}$$

$DataRecovered = (FeatureVector(p=1) \times TransformedData) + OriginalMean$

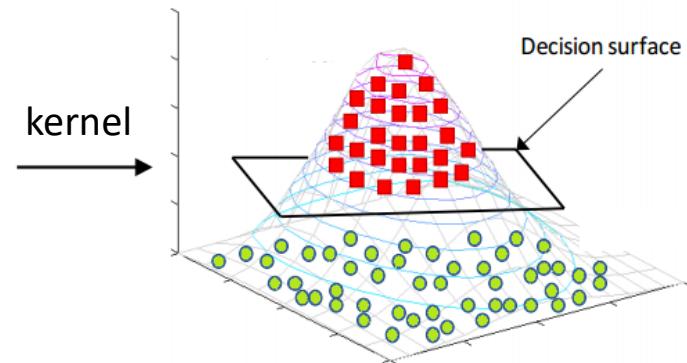
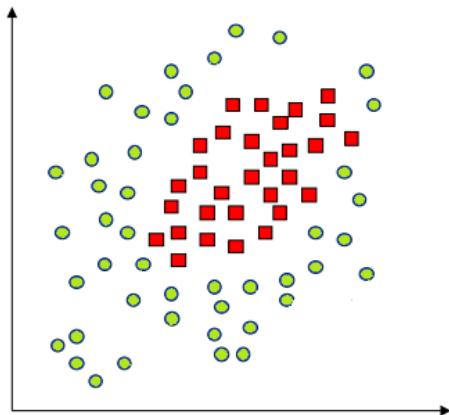
# Outline

- Dimensionality reduction: why and how
- Feature selection
- **Linear transformations**
  - algebra ground
  - principal component analysis
  - compression and reconstruction
  - **non-linear kernels**
  - **linear discriminant analysis**
- Subspace selection
- Evaluation
- Data reduction

# Nonlinearities using kernels



Linear projections will  
not detect this pattern



# Nonlinear PCA using kernels

- Traditional PCA applies linear transformation (ineffective for nonlinear data)
- **Solution:** apply nonlinear transformation to potentially higher-dimensional spaces

$$\varphi: \mathbf{x} \rightarrow \varphi(\mathbf{x})$$

- how? apply the kernel trick: PCA rewritten in terms of dot product

$$K(\mathbf{x}_i, \mathbf{x}_j) = \varphi(\mathbf{x}_i) \bullet \varphi(\mathbf{x}_j)$$

- **Example**  $\varphi: (a_{i1}, a_{i2}, a_{i3}) \rightarrow (a_{i1}, \sqrt{a_{i2}a_{i3}}, {a_{i1}}^2a_{i3})$

*simplified:* transform  $A = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$  accordingly and apply PCA

$$\mathbf{x}_i = (2, 4, 1) \rightarrow (2, 2, 4)$$

# Dimensionality reduction techniques

## Linear transformations

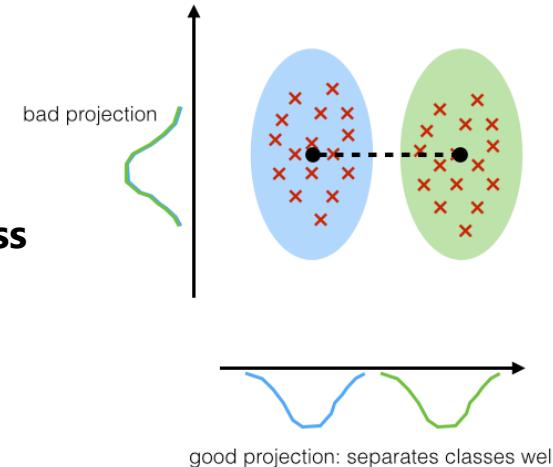
- **Eigenvalue analysis:** SVD, PCA and KL:  $O(nm^2)$  time
- **FastMap:** approximate searches in  $O(nm)$  time  
(optimal searches also in  $O(nm^2)$  time)
- **Multi-dimensional scaling:**  $O(nm^2)$  time
- **Random projections:**  $O(nm)$  time

## Non-linear transformations

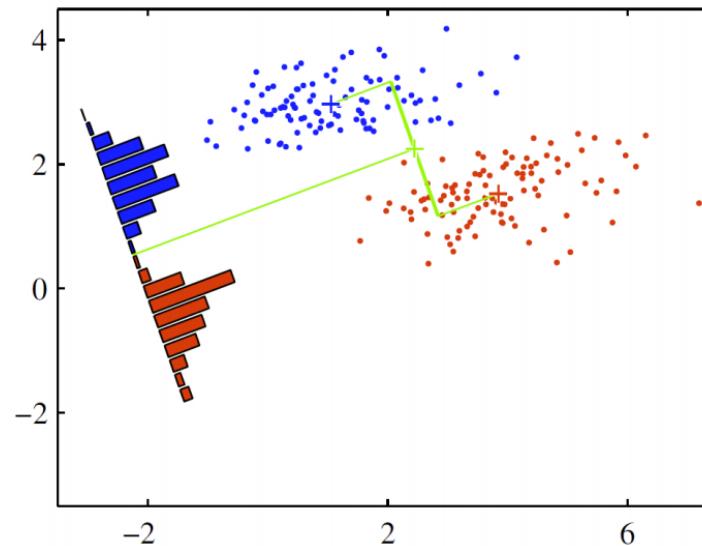
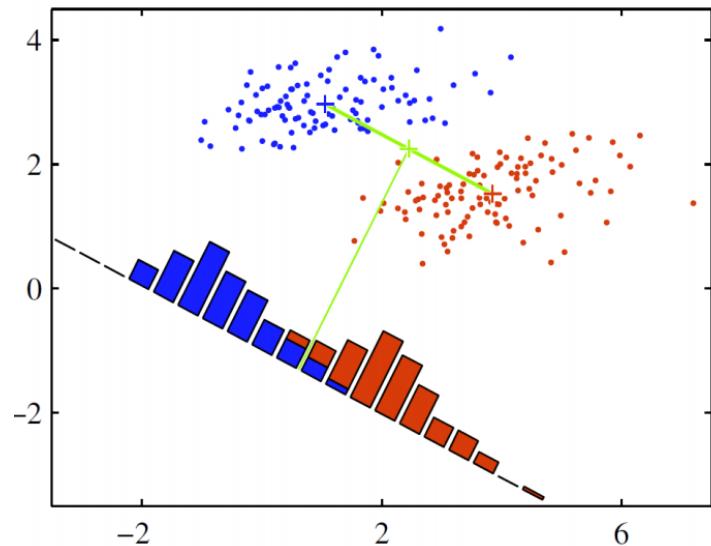
- **Kernel SVD/PCA**
- Locally linear embedding (**LLE**)
- Laplacian eigenmaps (**LEM**)
- Semidefinite embedding (**SDE**)

# Linear discriminant analysis (LDA)

- Challenges of PCA in supervised settings?
  - Reduction does not consider impact on the ability to discriminate output variables
- Goal: data transformation guaranteeing class separation
- Principle: pick a new dimension that
  - **maximize separation between projected classes**
  - **minimize variance of observations within each class**
- Solution: **LDA**
  - eigenvectors based on between-class and within-class covariance matrices



# Linear discriminant analysis (LDA)

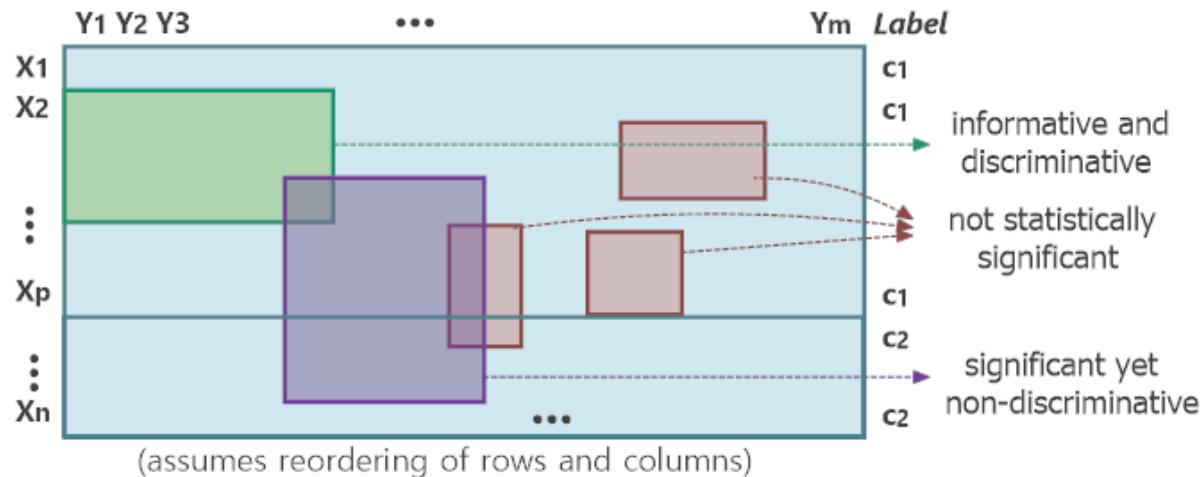


# Outline

- Dimensionality reduction: why and how
- Feature selection
- Linear transformations
  - algebra ground
  - principal component analysis
  - compression and reconstruction
  - non-linear kernels
  - linear discriminant analysis
- **Subspace selection**
- Evaluation
- Data reduction

# Subspace selection

- *Minimize overfitting:* remove uninformative regions (focus on informative/discriminative patterns only)
- *Minimize underfitting:* mine *all* relevant regions



# Subspace selection

- Addresses limitations of feature selection/extraction
  - single space → multiple compact spaces
  - goodness criteria computed from all observations  
→ goodness verified on a subset of observations
- Yet limited applicability
  - **pattern mining** and **biclustering**
  - **associative classification**
  - pattern-centric visualization

# Outline

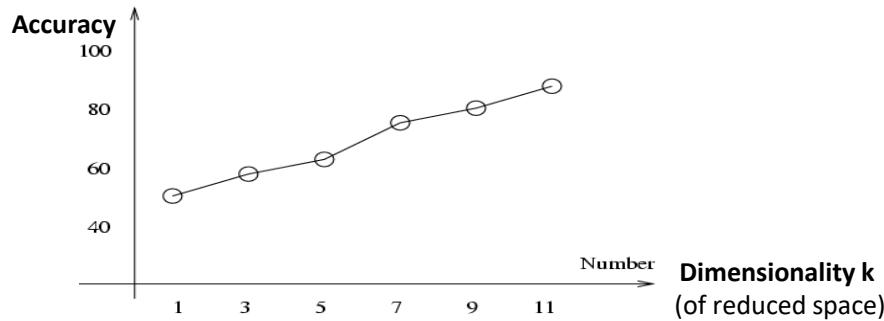
- Dimensionality reduction: why and how
- Feature selection
- Linear transformations
  - algebra ground
  - principal component analysis
  - compression and reconstruction
  - non-linear kernels
  - linear discriminant analysis
- Subspace selection
- **Evaluation**
- Data reduction

# Evaluation of low-dimensional spaces

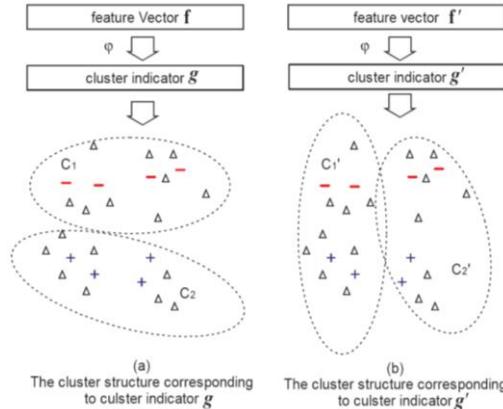
- Direct metrics
  - **match scores** against a reference set features/subspaces  
(only suitable for artificial data or data with prior knowledge)
- Indirect evaluation
  - **Learning on the original versus reduced data space**  
(e.g., predictive accuracy, goodness of resulting clusters)
    - before-and-after comparison
    - comparison using different learning algorithms
  - **Reconstruction error**
  - **Learning curves**

# Evaluation of low-dimensional spaces

- learning curves



- before-and-after fitness using clustering



# Outline

- Dimensionality reduction: why and how
- Feature selection
- Linear transformations
  - algebra ground
  - principal component analysis
  - compression and reconstruction
  - non-linear kernels
  - linear discriminant analysis
- Subspace selection
- Evaluation
- **Data reduction**

# Data reduction

Dimensionality reduction is one of the multiple forms of data reduction

- *data reduction* similarly considered to aid learning and time-memory efficiency

Other **data reduction strategies**:

- **domain reduction**: reduce data representation
  - *binning/discretization* of real-valued attributes, e.g. continuous variable onto ranges
  - *cardinality reduction* of nominal and ordinal attributes: aggregating categories
- **data size reduction (subsampling)**: reduce the number of instances
  1. *simple random subsample (SRS)*: randomly remove observations  
(either without or with replacement)
  2. *balanced sample*: remove observations guaranteeing that the reduced data satisfies a predefined criterion (e.g. balanced number of classes)

# Data reduction

## Subsampling (*continue*)

3. *stratified sample*: observations are divided into mutually disjointed parts, called strata, and removal (SRS) is done at each stratum
4. *data clustering*
  - clustering the data and use centroids instead of the actual data
5. *data condensation*
  - obtain minimal data set for correctly classifying all original observations
  - emerges from the fact that naive sampling (random or stratified sampling) is not suitable to learn from noisy data (dependent on algorithms)
6. *data squashing*
  - compress ("squash") data in such a way that a statistical analysis carried out on the compressed or original data yields the same outcome

# Thank you!

Rui Henriques

[rmch@tecnico.ulisboa.pt](mailto:rmch@tecnico.ulisboa.pt)